

Attested Transition Execution for DAZE-Style Simulator Backdoors

Jeong Woo Lee and Yongje Shin, Semyung University
jeongwoolee340@gmail.com, dudu jr@semyung.ac.kr
USENIX Security 2026 Poster Abstract

Abstract

RL agents are increasingly trained on third-party simulators, making the simulator a supply-chain trust anchor. DAZE shows that a malicious wrapper can implant reward-free backdoors by changing the learner-visible transition after an action is submitted, so action-level logging can miss the attack [1]. Prior RL backdoor work mainly studies policy triggers or sanitization [2, 3, 4]; we protect the simulator data plane. *Attested Transition Execution* (ATE) admits an update only when an approved measured step relation, exact replay, or enrolled receipt service vouches for the full transition record. On original DAZE Hopper/Reacher PPO, ATE restores clean-level ASR/return and admits zero tampered updates; SAC Reacher drops from ASR 0.9996 to 0.

Claim and Contributions

Claim. For an approved measured closure with an untrusted wrapper, DAZE-style transition tampering cannot feed forged updates into training without failing ATE admission. **Contributions.** (i) We frame wrapper backdoor defense as transition attestation, not anomaly detection. (ii) Action-only logging is insufficient: Hopper termination tampering preserves $a_{\text{sub}} = a_{\text{exec}}$ yet reaches ASR 0.982. (iii) ATE restores original-DAZE Hopper/Reacher to clean-level ASR/return and admits zero poisoned updates over 3.94M checks.

Design




Figure 1. ATE inserts a transition-admission gate between the wrapper and learner.

ATE checks learner-visible transitions rather than wrapper self-reports, accepting a transition only if the approved step relation validates:

$$(s, o, a, T_\theta(a, s, t, \text{rng}), s', o', r, d, \tau).$$

Here s, o denote state/observation; a the submitted action; T_θ the measured executed action; s', o' post-state/observation; r reward; d termination/truncation; and τ receipt metadata. T_θ covers approved actuator transforms; receipt mode follows measured-evidence principles [5].

Positive Learning Evidence

Setting	Lrn.	Clean	Attack	ATE	Admit
H/R DAZE	PPO	.0177/.000	.843/.984	.0177/.000	0
Reacher	SAC	.0000	.9996	.0000	0
DAZE					
CartPole edit	DQN	.415	.961	.415	0/80,350
Replay	DQN	.000	.908/.933	.000	0/all
relabel					

Table 1. Positive learning evidence across MuJoCo, Gymnasium, and custom replay settings.

Table 1 reports Attack Success Rate (ASR); promoted rows preserve clean-level utility and admit zero poisoned updates.

Systems Evidence

ATE also passed 168 MuJoCo promotion checks, 57 software-root protocol cases, 71 AMC-ATE stress cases, and 1.57M Brax/JAX GPU receipt records with zero accepted violations or false rejects. Code and result tables are public [6].

Boundary and Deployment

The approved closure covers simulator core, configuration/assets, RNG schedule, observation/termination code, reward inputs, and allowed actuator transforms; the wrapper is untrusted. ATE guarantees admission integrity for that closure, not safety of a maliciously approved closure. Current evidence targets exact CPU replay and software roots; remote clusters can bind receipts to TPM-style integrity verification [7, 8].

References

- [1] E. Rathbun et al. Beware Untrusted Simulators – Reward-Free Backdoor Attacks in RL. ICLR, 2026.
- [2] P. Kiourti et al. TrojDRL: Trojan Attacks on Deep RL Agents. DAC, 2020.
- [3] L. Wang et al. BACKDOORL: Backdoor Attack against Competitive RL. IJCAI, 2021.
- [4] J. Cui et al. BadRL: Sparse Targeted Backdoor Attack Against RL. AAAI, 2024.
- [5] H. Birkholz et al. RFC 9334: RATS Architecture. IETF, 2023.
- [6] J. W. Lee and Y. Shin. ATE code and result tables. GitHub repository, 2026.
- [7] E. Todorov et al. MuJoCo: A Physics Engine for Model-Based Control. IROS, 2012.
- [8] G. C. Fedorkow et al. RFC 9683: RIV with TPMs. IETF, 2024.